

Kivi - Product Vision

1. Market Context

The voice tech space has evolved significantly, moving well beyond basic dictation into AI-assisted speech. Tools like Wispr Flow, Superwhisper, and TalkTastic already allow people to talk naturally while clean-up algorithms handle punctuation and structure. By tapping into application context, these tools deliver far better writing outcomes. People clearly love speaking at the speed of thought to reduce friction, and newer offerings demonstrate a growing appetite for personalization through custom styles, modes, and vocabulary.

Even so, most products on the market still treat voice primarily as a method to **enter or improve text**. TalkTastic frames itself as a voice keyboard for typing or pasting into apps, and Superwhisper builds its modes around processing spoken input.

The infographic compares two AI voice tools, VoicelInk and Wispr Flow, using a central 'VS' graphic. It features icons for each tool, lists of features, and illustrations of laptops and coffee cups. A central text block encourages users to compare features to find the best tool for their workflow. A bottom call-to-action box prompts users to discover key differences.

VoicelInk VS Wispr Flow

Compare features, find the best AI voice tool for your workflow.

VoicelInk	Wispr Flow
✓ Works in Any App ✓ Real-time Dictation ✓ AI Formatting & Cleanup ✓ High Accuracy	✓ Automatic Formatting ✓ Punctuation & Cleanup ✓ Cross-platform Sync ✓ Built for Speed

VoicelInk	Wispr Flow
- Works in Any App - Real-time Dictation - AI Powered Accuracy - Privacy Focused	- Cross-platform Access - Auto Punctuation & Cleanup - Built for Speed - Cloud Sync

Speak naturally, write effortlessly.

Think it, Flow it.

Discover the **key differences** and find out which AI voice tool fits your workflow best. →

As these solutions get smarter, a fresh challenge emerges: finding the right balance between software interpreting intent and preserving exact user meaning. Unwanted rewriting, formatting bugs, reliability issues, and privacy concerns still frustrate users, while multilingual and code-mixed speech remain secondary to the core user experience.

This creates a major opportunity for voice to become more than just another way to type. It can be a seamless, natural way to interact with **what is already on the screen, across different applications, languages, and contexts.**

2. Changing the Relationship with Computers

For decades, computers have trained us on how they expect us to work.

We are trained where to click, which buttons to push, which keyboard shortcuts to memorize, and how to structure our commands so the system understands us.

Voice technology has largely followed this same pattern: speak something out loud, convert it to text, and take it from there.

Kivi is designed to completely redefine that relationship.

3. Kivi as a Voice Layer

Kivi works as an ever-present voice utility rather than another standalone application you have to open.

It seamlessly stays by your side whenever you use Chrome, Gmail, Docs, Cursor, ChatGPT, or any other tool.

You can bring Kivi Anywhere straight into whatever project you are working on, while Kivi Mic upgrades your existing microphone without disrupting your established workflow, much like Gboard.

Whatever app you are using stays right where it is.

Kivi steps in, gives you a hand, and quietly steps back out of the way.

4. The Core Interaction

The core interaction is so intuitive it can be summed up in a single sentence:

Start by pointing at something, talk through it, let Kivi understand it, preview the result, and finish by applying it.

Simply select something important on your screen and watch it highlight in Kivi green.

That selection tells Kivi exactly what you are focusing on.

From there, just say what you need, like **“shorten this”, “explain this”, “put this into bullet points”, or “make this sound less formal”**.

Your screen gives Kivi the location context.

Your voice gives Kivi the direction.

Kivi puts them together to figure out your exact intent.

5. Working With Information

Kivi shines when you are handling information.

It effortlessly converts a raw stream of consciousness into clear structure.

You can write and edit both copy and code fluidly without pausing to craft detailed prompts.

It helps you digest papers, articles, docs, and code bases, even reading explanations out loud via TTS while you read along.

Through Communicate, you can adapt a single core idea for different audiences and tones seamlessly.

The system saves key details you want to keep, helping Kivi learn what matters most to you over time.

6. Five Ways Kivi Helps

These capabilities come together in five simple ways Kivi supports your workflow:

Think

Think turns unorganized thoughts into concrete decisions, key questions, clear tasks, or structured outlines.

Write

Write assists you in drafting and editing text as well as code.

Understand

Understand breaks down complex material you are reading or viewing so it is easy to grasp.

Communicate

Communicate helps tailor the same underlying message across different people, languages, or situations, whether in voice or text.

Remember

Remember stores information you might otherwise lose, carries valuable details forward, and learns your personal style preferences.

7. Language as a Layer

Language should feel completely effortless.

You should never have to stop and consciously decide whether you are speaking Tamil or English.

Mixing languages is second nature to people, particularly in India, and Kivi embraces that reality.

Speech-to-text should capture real human conversation, while Kivi delivers output in whichever language and format best fits the context.

8. Kivi Mic and Speech-to-Speech



Kivi Mic elevates this experience even further.

Powered by Sarvam's speech-to-speech (STS) technology, language barriers disappear when communicating with others.

You can speak in Tamil while sending a voice note in Hindi, or hold a live conversation where the other person hears everything translated into their native language.

The ultimate goal is for people to express raw thoughts without translating them in their heads first.

People simply talk naturally, and Kivi takes care of the translation behind the scenes.

9. Personalization

Personalization lies at the heart of Kivi's future.

Kivi is designed not to be a generic, one-size-fits-all assistant.

You can easily teach Kivi uncommon words, set up custom voice shortcuts, create new modes from just a couple of examples, lock in preferred writing styles, and correct mistakes on the fly.

Over time, Kivi deepens its understanding of your unique vocabulary, daily routines, and shorthand habits, far beyond just recognizing your voice.

(Personalization Moat)

That personalized adaptation is what makes Kivi irreplaceable.

Other platforms might offer dictation, text-to-speech, translation, or rewriting features.

However, no competitor can instantly replicate your custom dictionary, shortcuts, tone preferences, custom modes, and subtle nuances that Kivi can build with you over time.

10. Human in the Loop

At the same time, Kivi understands its boundaries.

When ambiguity or any decision making arises, it presents its interpretation for confirmation before executing any changes on your behalf.

Human in the Loop (HITL)

The vision is simple:

Express what you want naturally using your voice, and let your computer effortlessly grasp both your intent and full screen context.